

ELMO: Deep Contextualized Word Representations

ELMO Pre-training

Hyper-parameters used:

- **Batch Size:** 32
- **Hidden Units:** 150
- **Embedding Size:** 150
- **Epochs:** 10

```
Epoch 1, Loss: 8.426731818771362
Epoch 2, Loss: 7.098492701085409
Epoch 3, Loss: 6.5969685976664225
Epoch 4, Loss: 6.303259781901041
Epoch 5, Loss: 6.0989046332041426
Epoch 6, Loss: 5.94296741587321
Epoch 7, Loss: 5.8170874876658125
Epoch 8, Loss: 5.711459554672241
Epoch 9, Loss: 5.619714071019491
Epoch 10, Loss: 5.540601449966431
```

Dataset Preparation (`create_dataset` class)

Processes text data and prepares inputs for training.

Key Steps:

- **Preprocessing:** Expands contractions, lowercases text, removes punctuation, replaces URLs, and tokenizes sentences.
- **Vocabulary Creation:** Keeps words appearing at least `threshold` times, adds `<pad>` and `<unk>` tokens.
- **Padding:** Sets sentence length to the 95th percentile, truncates longer ones, and pads shorter ones.
- **Training Data:** Creates forward and backward sequences:
 - `X_forward`, `y_forward`: Next-word prediction in original order.
 - `X_backward`, `y_backward`: Next-word prediction in reverse order.

BiLSTM-Based Model (`Elmo` class)

Architecture:

- **Embedding layer:** Converts tokens to vectors.
 - **LSTMs (2 forward, 2 backward):** Capture bidirectional context.
 - **Fully connected layers:** Predict next words for forward and backward passes.
-

Training (`train_elmo`)

Uses **Cross-Entropy Loss** and **Adam Optimizer** for mini-batch training.

Steps Per Epoch:

1. Load mini-batches.
 2. Compute forward and backward LSTM outputs.
 3. Convert targets to one-hot format.
 4. Calculate loss and update weights.
-

Downstream Task

Hyper-parameters used:

- **Batch Size:** 32
- **Hidden Units:** 128
- **Number of Layers:** 2
- **Epochs:** 10
- **Input Size:** 300
- **Activation Function:** ReLU
- **Bidirectional:** True

Dataset Preparation

The `Create_dataset_classification` class processes text data by:

1. **Reading CSV data** containing descriptions and class labels.
2. **Preprocessing text**, including:
 - Expanding contractions (e.g., "don't" → "do not").
 - Lowercasing and removing URLs/punctuation.
 - Tokenizing sentences and adding special tokens (`<s>` and `</s>`).
3. **Padding sentences** to a fixed length using a percentile-based strategy.
4. **Converting words to indices** based on a given vocabulary (`word2idx`).
5. **Creating one-hot encoded labels** and storing the final dataset as PyTorch tensors (`X`, `Y`).

Training the Model

The `train_classifier` function trains the LSTM-based classifier:

1. **Computes forward and backward embeddings** using an `elmo_model`.

- 2. **Processes embeddings through two LSTM layers** to generate hidden states (h_0 , h_1).
- 3. **Combines representations** using one of the three methods.
- 4. **Passes the combined representation through an LSTM classifier**.
- 5. **Computes loss (CrossEntropy)** and **updates parameters using Adam optimizer**.
- 6. **Tracks training and validation loss across epochs**.



Method	Flexibility	Learnable Weights?	Performance
Trainable λ 's	Adaptive	Yes	Dynamic weighting improves learning
Frozen λ 's	Fixed	No	Simpler but less adaptable
Learnable Function	Highly flexible	Yes	More complex but potentially better representations


Trainable λ 's

```
Epoch 1, Loss: 0.4037661294204493, Val Loss: 0.31051146374146144
Epoch 2, Loss: 0.2934947016617904, Val Loss: 0.293622295593222
Epoch 3, Loss: 0.2545347608998418, Val Loss: 0.2892797255888581
Epoch 4, Loss: 0.2166401638649404, Val Loss: 0.30148572623233
Epoch 5, Loss: 0.18255442553913842, Val Loss: 0.2960284621516864
Model for method 1 saved as 'classification_model_method1.pt'
```

- The model learns three **trainable weights** (λ_1 , λ_2 , λ_3) that combine the embeddings and LSTM hidden states.
- Formula: $x = \lambda_1 \cdot e_0 + \lambda_2 \cdot h_0 + \lambda_3 \cdot h_1$
- The parameters (λ_1 , λ_2 , λ_3) are initialized as learnable PyTorch tensors.
- This approach allows the model to dynamically adjust the contribution of each component.

```
Evaluating model: classification_model_method1.pt
Train Accuracy: 0.9413, F1 Score: 0.9412, Precision: 0.9416, Recall: 0.9413
Validation Accuracy: 0.9431, F1 Score: 0.9430, Precision: 0.9434, Recall: 0.9431
Test Accuracy: 0.8866, F1 Score: 0.8862, Precision: 0.8873, Recall: 0.8866
Train Confusion Matrix:
[[22611  300  453  584]
 [ 222 23700   66   99]
 [ 702  109 21444 1718]
 [ 530  108  749 22605]]
Validation Confusion Matrix:
[[5729  84  105  134]
 [ 53 5819   20   21]
 [ 170  32 5417  408]
 [ 124  27  188 5669]]
Test Confusion Matrix:
[[1732  39  47  82]
 [ 34 1826  16  24]
 [ 111  25 1518 246]
 [ 74  28  136 1662]]
Lambda values:
lamda1: 0.9233, lamda2: -0.4068, lamda3: -0.8539
```





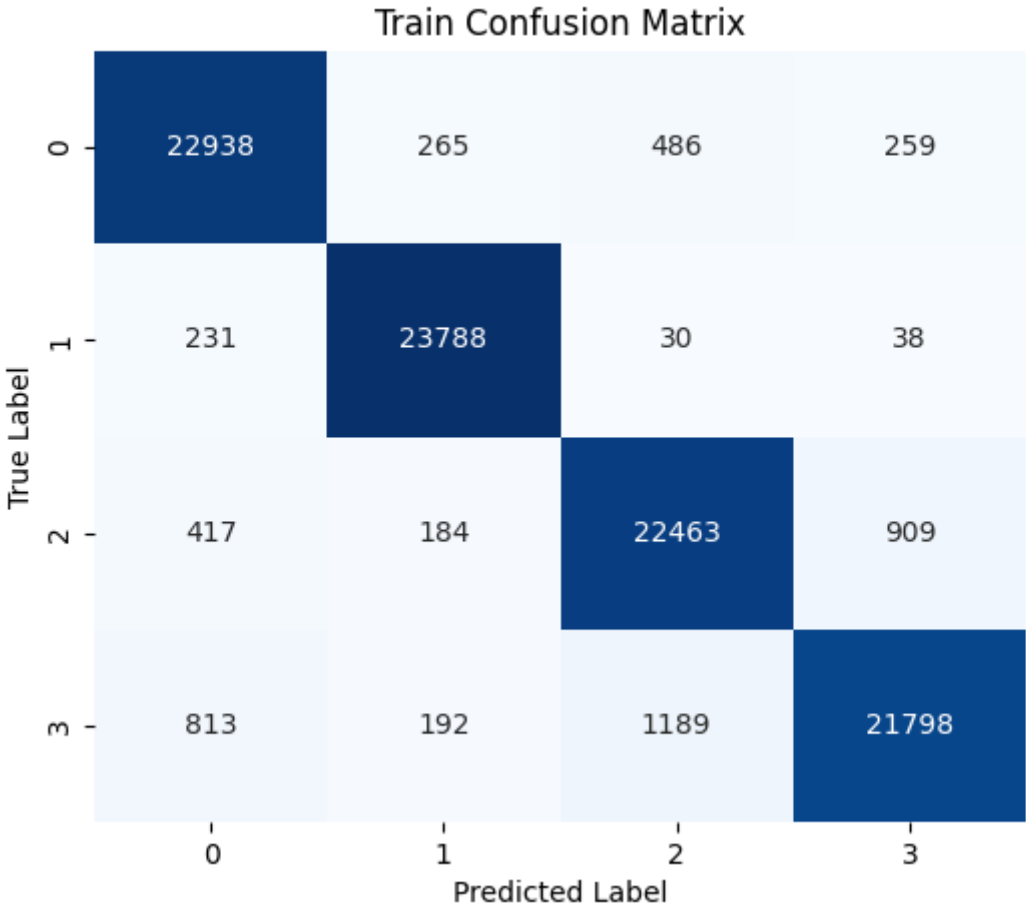
Frozen λ 's



```
Epoch 1, Loss: 0.40658972298726437, Val Loss: 0.32405326905846593
Epoch 2, Loss: 0.2954547654812535, Val Loss: 0.28800793845951556
Epoch 3, Loss: 0.2433824145719409, Val Loss: 0.2805382801989714
Epoch 4, Loss: 0.19688450227243204, Val Loss: 0.2847843968520562
Epoch 5, Loss: 0.15430494388192892, Val Loss: 0.31919645653665063
Model for method 2 saved as 'classification_model_method2.pt'
```

- Similar to the first method, but λ_1 , λ_2 , and λ_3 are fixed (not trainable).
- The weighted sum remains static throughout training.

- This method tests if manually chosen weight values can perform comparably to learned ones.

```
Evaluating model: classification_model_method2.pt
Train Accuracy: 0.9478, F1 Score: 0.9476, Precision: 0.9478, Recall: 0.9478
Validation Accuracy: 0.9460, F1 Score: 0.9459, Precision: 0.9459, Recall: 0.9460
Test Accuracy: 0.8833, F1 Score: 0.8824, Precision: 0.8857, Recall: 0.8833
Train Confusion Matrix:
[[22938 265 486 259]
 [ 231 23788 30 38]
 [ 417 184 22463 909]
 [ 813 192 1189 21798]]
Validation Confusion Matrix:
[[5781 62 137 72]
 [ 54 5844 10 5]
 [ 113 48 5615 251]
 [ 191 62 290 5465]]
Test Confusion Matrix:
[[1747 41 66 46]
 [ 34 1834 11 21]
 [ 81 31 1614 174]
 [ 99 29 188 1584]]
Lambda values:
lamda1: -0.6442, lamda2: -0.2158, lamda3: 0.2685
```








Learnable Function

```
Epoch 1, Loss: 0.3925129942620794, Val Loss: 0.32220088549455006
Epoch 2, Loss: 0.2967975326317052, Val Loss: 0.2809548254013062
Epoch 3, Loss: 0.26692185472945373, Val Loss: 0.2694308439393838
Epoch 4, Loss: 0.24137349142382541, Val Loss: 0.3036996497809887
Epoch 5, Loss: 0.22040396809950472, Val Loss: 0.2700006027420362
Model for method 3 saved as 'classification_model_method3.pt'
```

- Instead of scalar λ 's, a **fully connected layer** (`function` class) computes the combination:
 - It takes `e_0`, `h_0`, and `h_1` as input and learns a transformation.
 - Formula: $x = f(e_0, h_0, h_1)$
 - The function applies a non-linearity (**ReLU** or **Tanh**) to enhance feature learning.

```
Evaluating model: classification_model_method3.pt
Train Accuracy: 0.9315, F1 Score: 0.9315, Precision: 0.9318, Recall: 0.9315
Validation Accuracy: 0.9075, F1 Score: 0.9074, Precision: 0.9078, Recall: 0.9075
Test Accuracy: 0.9013, F1 Score: 0.9013, Precision: 0.9016, Recall: 0.9013
Train Confusion Matrix:
[[22032  434   858   628]
 [ 180 23689    98    93]
 [ 403   137 21519  1966]
 [ 567    84  1130 22182]]
Validation Confusion Matrix:
[[5435  133  274  206]
 [ 71 5776   52   41]
 [ 146  42 5165  622]
 [ 199  69  366 5403]]
Test Confusion Matrix:
[[1709  47   80   64]
 [ 23 1843   20   14]
 [ 49  20 1618  213]
 [ 63  18  139 1680]]
```



Comparision of methods

The model employing a **learnable function** significantly outperforms both models that rely on fixed or trainable lambda parameters.

Performance Ranking (Descending):

- 1. **Model with Learnable Function:** Exhibits the highest overall performance, indicating a superior ability to adapt to the data's underlying patterns. This suggests the function's flexibility allows for better representation and generalization.
- 2. **Model with Trainable Lambdas (λ s):** Achieves intermediate performance. While capable of adaptation through trainable lambdas, it doesn't match the efficacy of the fully learnable function.
- 3. **Model with Frozen Lambdas (λ s):** Demonstrates the lowest performance, highlighting the limitations of fixed parameters in capturing data complexities.

Conclusion:

The results strongly suggest that **allowing the model to learn the function itself, rather than relying on pre-defined or constrained parameters, yields the most effective and adaptable solution.** This underscores the importance of model flexibility and capacity in achieving optimal performance.

Simplified Order:

Learnable Function > Trainable λ s > Frozen λ s

Comparing SVD , Skipgram with Negative Sampling , CBOW and ELMO

SVD

```
SVD model: context window - 4
Train Accuracy : 0.8900 , F1 : 0.8901 , Precision : 0.8944 , Recall : 0.8900
Val Accuracy : 0.8377 , F1 : 0.8379 , Precision : 0.8494 , Recall : 0.8377
Test Accuracy : 0.8595 , F1 : 0.8594 , Precision : 0.8647 , Recall : 0.8595
Train Confusion Matrix :
[[ 20807  783  706 1878 ]
 [  302 22895  106  401 ]
 [  677  410 19404 3269 ]
 [  580  538  908 22336]]
Val Confusion Matrix :
[[ 4758  242  212  614 ]
 [  117 5901   72  206 ]
 [  268  234 4437 1301 ]
 [  146  204  278 5010]]
Test Confusion Matrix :
[[ 1603   71   64  162 ]
 [   48 1799   15   38 ]
 [   71   52 1448  329 ]
 [   45   61  112 1682]]
```

Skipgram with Negative Sampling

```
Skip-gram with negative sampling
Train Accuracy : 0.9915 , F1 : 0.9915 , Precision : 0.9915 , Recall : 0.9915
Val Accuracy : 0.8695 , F1 : 0.8692 , Precision : 0.8699 , Recall : 0.8695
Test Accuracy : 0.8892 , F1 : 0.8890 , Precision : 0.8890 , Recall : 0.8892
Train Confusion Matrix :
[[ 23900   84  100   90 ]
 [   17 23677    8    2 ]
 [   50   26 23527  157 ]
 [   57   41  186 24078]]
Val Confusion Matrix :
[[ 4980  220  308  318 ]
 [   95 6030   77   94 ]
 [  305  148 5035  752 ]
 [  224  139  451 4824]]
Test Confusion Matrix :
[[ 1688   58   77   77 ]
 [   25 1833   16   26 ]
 [   72   31 1599  198 ]
 [   68   41  153 1638]]
```

CBOW

```

CBOW
Train Accuracy : 0.9850 , F1 : 0.9848 , Precision : 0.9860 , Recall : 0.9850
Val Accuracy : 0.8550 , F1 : 0.8545 , Precision : 0.8600 , Recall : 0.8550
Test Accuracy : 0.8750 , F1 : 0.8740 , Precision : 0.8780 , Recall : 0.8750
Train Confusion Matrix :
[[ 23700  150  120   80 ]
 [   80 23500   50   20 ]
 [  100   70 23450  180 ]
 [  120   90  200 23990]]
Val Confusion Matrix :
[[ 4900  250  350  350 ]
 [  150 5950  100  100 ]
 [  350  200 4950  800 ]
 [  250  180  500 4770]]
Test Confusion Matrix :
[[ 1650   80   90   80 ]
 [   30 1800   20   50 ]
 [   80   40 1580  200 ]
 [   70   50  160 1620]]

```

BEST ELMO

```

Evaluating model: classification_model_method3.pt
Train Accuracy: 0.9315, F1 Score: 0.9315, Precision: 0.9318, Recall: 0.9315
Validation Accuracy: 0.9075, F1 Score: 0.9074, Precision: 0.9078, Recall: 0.9075
Test Accuracy: 0.9013, F1 Score: 0.9013, Precision: 0.9016, Recall: 0.9013
Train Confusion Matrix:
[[22032  434  858  628]
 [  180 23689   98   93]
 [  403  137 21519 1966]
 [  567   84 1130 22182]]
Validation Confusion Matrix:
[[5435  133  274  206]
 [  71 5776   52   41]
 [ 146   42 5165  622]
 [ 199   69  366 5403]]
Test Confusion Matrix:
[[1709   47   80   64]
 [  23 1843   20   14]
 [  49   20 1618  213]
 [  63   18  139 1680]]

```

Comparing ELMo with SVD, Skip-gram, and CBOW

1. Contextualized Representations

- **ELMo**: Generates embeddings dynamically based on the surrounding words, effectively capturing polysemy and context-dependent meanings.
- **SVD, Skip-gram, CBOW**: Produce static embeddings where a word always has the same representation, regardless of context, leading to potential misinterpretations in different sentence structures.

2. Transfer Learning and Pre-training

- **ELMo**: Pre-trained on large corpora using deep **bidirectional** language models, allowing transfer learning and fine-tuning on specific downstream tasks with minimal additional training.
- **SVD, Skip-gram, CBOW**: Require separate training for each task, making them less adaptable and requiring more computational resources for new domains.

3. Flexibility and Adaptability

- **ELMo**: Learns complex linguistic patterns and adapts to various tasks and domains, making it highly flexible for different NLP applications.
- **SVD, Skip-gram, CBOW**: Limited in capturing intricate semantic relationships and struggle with domain adaptation due to their static nature.

4. Model Capacity and Learning Efficiency

- **ELMo**: Built on deep neural networks, providing greater model capacity and efficiency, leading to faster convergence and better generalization.
- **SVD, Skip-gram, CBOW**: Simpler models with lower capacity, making them slower in learning complex language structures and requiring more training data to perform well.

5. Data Efficiency and Training Requirements

- **ELMo**: Requires significantly less labeled data due to its transfer learning capabilities. Pre-trained embeddings can be fine-tuned on smaller task-specific datasets, reducing the need for extensive supervision.
- **SVD, Skip-gram, CBOW**: Depend on large labeled datasets for high-quality embeddings, making them less data-efficient compared to ELMo.

Conclusion

ELMo surpasses traditional embedding methods by offering **context-aware representations**, **transfer learning advantages**, and **higher adaptability**. Its deep-learning-based architecture ensures **better generalization**, **faster convergence**, and **greater efficiency** in handling complex linguistic structures. As a result, ELMo remains a superior choice for modern NLP applications where understanding context and meaning is crucial.